

Binary Search Trees

Contents:

- What is a binary search tree?
- Querying a binary search tree
- Insertion and deletion
- Randomly built binary search trees (*)

Search Trees

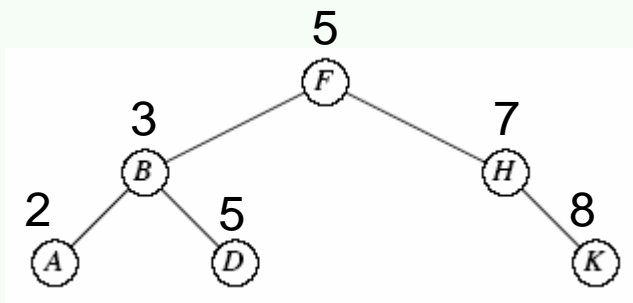
- Data structures that support many dynamic-set operations
 - SEARCH, MINIMUM, MAXIMUM, PREDECESSOR, SUCCESSOR, INSERT, DELETE
- Can be used as both a dictionary and as a priority queue
- Basic operations take time proportional to the height of the tree
- For complete binary tree with n nodes: worst case $\Theta(\lg n)$
- For linear chain of n nodes: worst case $\Theta(n)$
- Different types of search trees include binary search trees, red-black trees (Ch13), and B-trees (Ch18)
- We will cover binary search trees, tree walks, and operations on binary search trees

Binary Search Trees

Important data structure for dynamic sets

- Accomplish many dynamic-set operation in $O(h)$ time
- Use a linked data structure to represent a binary tree
 - Node : object contains **key**, **left**, **right**, **p**
 - root[T]** : root of tree T, $p[\text{root}[T]] = \text{NIL}$
- Binary-search-tree property
 - If y is in **left** subtree of x , then $\text{key}[y] \leq \text{key}[x]$
 - If y is in **right** subtree of x , then $\text{key}[y] \geq \text{key}[x]$

Height of a tree



Inorder-Tree-Walk

Print keys in a binary search tree in order, recursively

- Check to make sure that x is not NIL
- Recursively, print the keys of the nodes in x 's left subtree
- Print x 's key
- Recursively, print the keys of the nodes in x 's right subtree

INORDER-TREE-WALK(x)

if $x \neq \text{NIL}$

then INORDER-TREE-WALK($\text{left}[x]$)

print $\text{key}[x]$

INORDER-TREE-WALK($\text{right}[x]$)

e.g. ABDFHK

Correctness: by induction directly from the binary-search-tree property

Time: $O(n)$ time for a tree with n nodes,
because we visit and print each node once

Searching a Binary Search Tree

- Search for the element with key= k ,
 $\text{TREE-SEARCH}(\text{root}[T], k)$

TREE-SEARCH(x, k)

if $x = \text{NIL}$ or $k = \text{key}[x]$

then return x

if $k < \text{key}[x]$

then return $\text{TREE-SEARCH}(\text{left}[x], k)$

else return $\text{TREE-SEARCH}(\text{right}[x], k)$

e.g. search for D

- **Time:** The algorithm recurses, visiting nodes on a downward path from the root. Thus, running time is $O(h)$

Minimum and Maximum

The binary-search-tree property guarantees that

- The **minimum** (**maximum**) key of a binary search tree is located at the **leftmost** (**rightmost**) node

Traverse the appropriate pointers (*left* or *right*) until **NIL** is reached.

```
TREE-MINIMUM(x)
  while left[x] ≠ NIL
    do x ← left[x]
  return x
```

```
TREE-MAXIMUM(x)
  while right[x] ≠ NIL
    do x ← right[x]
  return x
```

Time: Both procedures visit nodes that form a **downward** path from the root to a leaf. Both procedures run in **$O(h)$** time

Successor and Predecessor

- Assuming that all keys are **distinct**,
 - $y = \text{successor}[x] \rightarrow \text{key}[y]$ is the **smallest key** $> \text{key}[x]$
 - $y = \text{predecessor}[x] \rightarrow \text{key}[y]$ is the **largest key** $< \text{key}[x]$
- We can find x 's successor based entirely on the tree structure. **No** key comparisons are necessary
- If x has the largest key in the binary search tree, then we say that x 's successor is **NIL**

To identify y , the successor of x , there are two cases:

- If node x has a **non-empty right subtree**, then y is the **minimum** in x 's right subtree (i.e. the **leftmost** node in the **right** subtree)
- If node x has an **empty right subtree**,
 - x **must NOT** be in y 's **right** subtree, and be the **maximum** in y 's **left** subtree
 - y is the **lowest ancestor** of x whose **left child** is also an **ancestor** of x
 - i.e. if we move up from x , y is the **first** ancestor we encounter when we go **right**

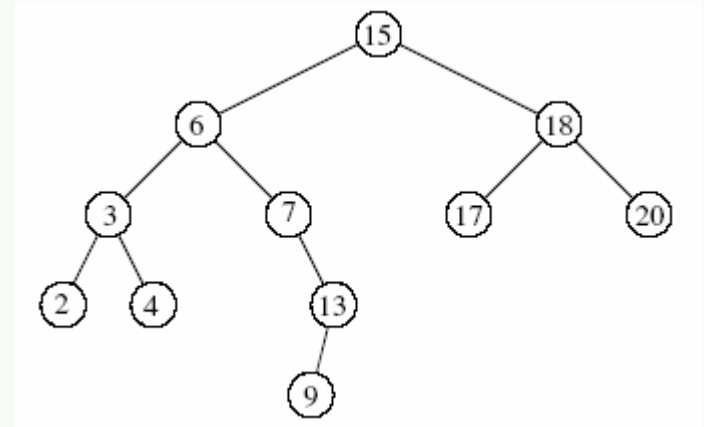
Searching for Successor

TREE-SUCCESSOR(*x*)

```

if right[x]  $\neq$  NIL
then return TREE-MINIMUM(right[x])
y  $\leftarrow$  p[x]
while y  $\neq$  NIL and x = right[y]
    do x  $\leftarrow$  y
       y  $\leftarrow$  p[y]
return y

```



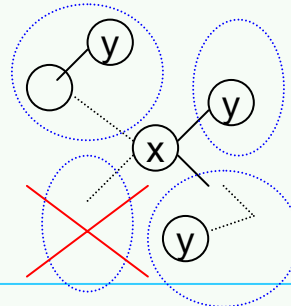
Q: the successor of the node with key value 15,6,4,17=? A: 17,7,6,18

Q: the predecessor of the node with key value 15,6,4,17=? A: 13,4,3,15

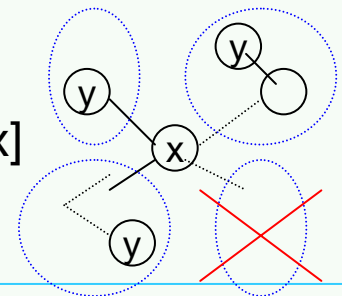
- TREE-PREDECESSOR is symmetric to TREE-SUCCESSOR
i.e. change *right*[*x*] by *left*[*x*] & TREE-MAXIMUM by TREE-MINIMUM

Time: $O(h)$

$y = \text{successor}[x]$



$y = \text{predessor}[x]$

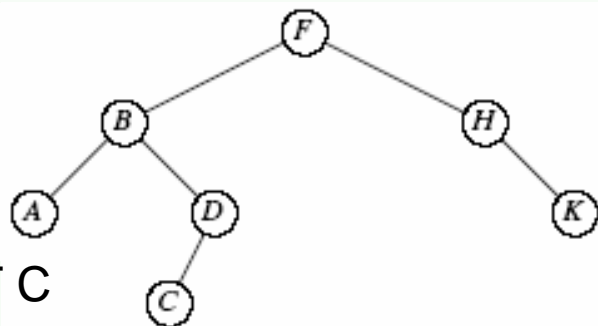


Insertion

- To insert z with $key[z] = v$, $left[z] = NIL$, and $right[z] = NIL$

Idea: for a node x , compare $key[z]$ and $key[x]$

- and $y = p[x]$,
- If $key[z] < key[x]$, z should be in x 's left subtree
else z is in x 's right subtree
- Diving:** Record $y = x$,
dive on the left (or right) subtree of x
by $x = left[x]$ (or $x = right[x]$), thus $y = p[x]$
- As long as $x \neq NIL$, we keep diving as above
- Then compare $key[z]$ and $key[y]$
- If $key[z] < key[y]$, then
- else $right[y] = z$



Time: $O(h)$

e.g. INSERT C

TREE-INSERT(T, z)

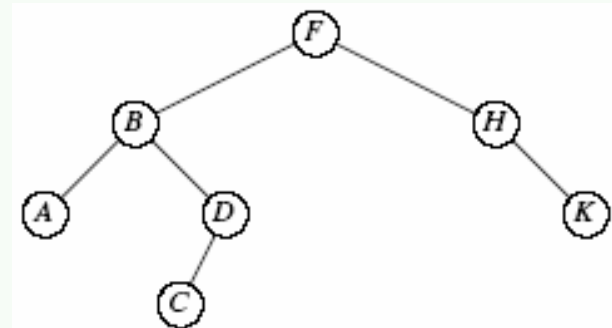
```

 $y \leftarrow NIL$ 
 $x \leftarrow root[T]$ 
while  $x \neq NIL$ 
  do  $y \leftarrow x$ 
    if  $key[z] < key[x]$ 
    then  $x \leftarrow left[x]$ 
    else  $x \leftarrow right[x]$ 
 $p[z] \leftarrow y$ 
if  $y = NIL$  //  $T$  was empty
then  $root[T] \leftarrow z$ 
else if  $key[z] < key[y]$ 
then  $left[y] \leftarrow z$ 
else  $right[y] \leftarrow z$ 
  
```

Deletion

TREE-DELETE(T, z) deletes node z from T

- **Case 1:** z has **no** children
 - Delete z by making the parent of z point to NIL, instead of to z
- **Case 2:** z has **one** child
 - Delete z by making the parent of z point to z 's child, instead of to z
- **Case 3:** z has **two** children
 - $y = \text{successor}[z]$ **must be** in z 's **right** subtree, and **have** either **no** children or **one** child.
 (y is the minimum node—with **no left** child—in z 's right subtree.)
 - Delete y from the tree (via Case 1 or 2).
 - Replace z 's key and satellite data with y 's.
 e.g. Case 1: delete K
 Case 2: delete H
 Case 3: delete B, swap it with C



TREE-DELETE(T, z)

//Determine which node y to splice out: either z or z 's successor.

if $left[z] = \text{NIL}$ or $right[z] = \text{NIL}$ // z has one or zero child

then $y \leftarrow z$

else $y \leftarrow \text{TREE-SUCCESSOR}(z)$ // z has two children

// x is set to a non-NIL child of y , or to NIL if y has no children.

if $left[y] \neq \text{NIL}$

then $x \leftarrow left[y]$

else $x \leftarrow right[y]$

// y is removed from the tree by manipulating pointers of $p[y]$ and x .

if $x \neq \text{NIL}$ // y has 1 child

then $p[x] \leftarrow p[y]$

if $p[y] = \text{NIL}$ // y is root

then $root[T] \leftarrow x$

else if $y = left[p[y]]$ // y is not root, splice out y

 then $left[p[y]] \leftarrow x$

 else $right[p[y]] \leftarrow x$

// If it was z 's successor that was spliced out, copy its data into z .

if $y \neq z$

then $key[z] \leftarrow key[y]$

 copy y 's satellite data into z

return y

Time: $O(h)$

Randomly built Binary Search Trees

- Given a set of n distinct keys. Insert them in random order into an initially empty binary search tree
- Each of the $n!$ permutations is equally likely
 - Different from assuming that every binary search tree on n keys is equally likely
 - the expected height of a randomly built binary search tree is $O(\lg n)$